

HERA-PdM: Lessons from Deploying Uncertainty-Gated Edge AI for Predictive Maintenance, from Bearings to a Truck Fleet

Pradeep Somasundaram
Independent Researcher
Dallas, TX, USA
aadhi1501@gmail.com

Abstract—Deploying predictive-maintenance (PdM) AI requires deciding where inference runs, how far predictions can be trusted and when humans act. We study HERA-PdM, in which cross-validated conformalized quantile regression (CQR-CV+) gives remaining-useful-life (RUL) intervals, the edge escalates only when its interval does not fix a unique maintenance action, and high-impact actions are SHAP-supported recommendations for human approval. On 17 FEMTO and 15 XJTU-SY run-to-failure bearings (bearing-grouped cross-validation, three seeds, tuned baselines), CQR-CV+ raised coverage to 0.94–0.96 and 0.93–0.99 (nominal 0.90). Versus point estimates with fixed thresholds, calibrated decisions cut critical snapshots left without action from 49.5% to 0.6% (FEMTO) and 34.2% to 9.2% (XJTU-SY), at a false-maintenance cost. On gradually degrading XJTU-SY bearings, escalating 79.1% of snapshots matched always-cloud safety with less false maintenance (53.2% vs. 69.1%), a gain not robust across bearings; on abrupt FEMTO failures it offered no benefit. The decision-sufficiency gate matched an interval-width gate, and per-bearing alarms were premature because a dozen calibration units are too few for unit-level risk control. On 23,550 Scania trucks, conformal risk control met its lateness target (9.9% vs. 40.5% for a fixed threshold) and HERA escalated 73% of readouts at always-cloud safety, although false alarms remained high.

Index Terms—AI deployment, predictive maintenance, edge AI, conformal prediction, explainable AI, human-in-the-loop, Digital Twin

I. INTRODUCTION

Predictive maintenance (PdM) estimates the remaining useful life (RUL) of machinery from condition data so that interventions are scheduled before failure [1]. Deploying it on a factory floor raises three questions that accuracy benchmarks ignore: *where* inference runs, since vibration is sensed at the machine and edge–cloud architectures keep routine inference local [2]; *how far* a prediction can be trusted, which calls for calibrated uncertainty rather than point estimates [3]; and *who* acts on it, since Industry 5.0 keeps humans in charge of consequential decisions [4].

Hierarchical PdM systems that combine edge devices, Digital Twins (DTs) and agents already exist [5], [6]; they escalate from edge to cloud on anomaly thresholds [5] or uncalibrated confidence [6]. Calibrated conformal deferral with guarantees has been developed for classification cascades [7] but not for RUL-based maintenance. We therefore ask *when* conformally calibrated RUL intervals computed on the edge are a useful

escalation signal, and what they cost in safety, false maintenance, alarm timing, communication and latency.

Key findings. On two public bearing datasets with contrasting degradation behaviour, (i) calibration removes most unsafe decisions of point-threshold maintenance; (ii) uncertainty-gated escalation pays off only when degradation is observable early; (iii) the precise gating rule matters less than the interval width; (iv) snapshot-level guarantees do not yield timely bearing-level alarms; and (v) on a fleet of 23,550 trucks, a unit-level lateness guarantee holds as calibrated (9.9% late at a 10% target versus 40.5% for a fixed threshold) and HERA sends 73% of readouts to the cloud, but false alarms remain high.

Contributions:

- 1) **HERA-PdM**, a reproducible four-layer framework with a *decision-sufficiency* gate that escalates only when the edge RUL interval straddles a maintenance-decision boundary, bounding locally resolved unsafe decisions by the edge miscoverage level, and CQR-CV+ calibration on all non-test assets.
- 2) A two-dataset, bearing-grouped benchmark of ten tuned quantile models with and without DT context, showing that CQR-CV+ restores nominal coverage and that larger fog/cloud models help only under gradual degradation.
- 3) Snapshot-, unit- and fleet-level evidence, with bootstrap intervals and a feasibility bound for unit-level risk control (Proposition 1) validated on 23,550 trucks, that turns findings (i)–(v) into deployment lessons.

II. RELATED WORK AND RESEARCH GAP

RUL prediction. Deep sequence models such as temporal convolutional networks [8] are standard RUL estimators; reviews [1] treat health-stage division, i.e. detecting when degradation becomes observable, as a prerequisite for prognosis, a property our two datasets [9], [10] contrast sharply. **Hierarchical PdM.** Ren et al. [2] run a lightweight TCN at the edge and refine RUL in the cloud; HAMA [5] escalates from a z-score edge pre-filter to fog ensembles with SHAP and generated operator text; EC-Distill-ZeroDiag [6] offloads uncertain fault-diagnosis cases from an edge agent to a cloud model. **Calibrated uncertainty.** Conformalized quantile regression (CQR) [3] and CV+ [11] give distribution-free coverage;

TABLE I
CLOSEST RELATED WORK (CODED FROM FULL TEXT, ABSTRACTS OR
PUBLIC CODE). ✓: PRESENT, ○: PARTIAL, -: ABSENT.

Work	DT	Edge	Agents	RUL	Calib.	UQ	Routing	XAI	HIL
Ren et al. [2]	-	✓	-	✓	-	-	-	-	-
HAMA [5]	-	✓	✓	○	-	-	○ anomaly	✓	○
EC-Distill [6]	-	✓	✓	-	-	-	✓ conf.	○	-
CAB cascade [7]	-	✓	-	-	✓	-	✓	-	-
Wang et al. [12]	-	-	-	✓	✓	-	-	-	-
HERA-PdM	○	✓	✓	✓	✓	-	✓	✓	✓




Fig. 1. HERA-PdM. The edge agent computes a calibrated RUL interval I_t^e ; the gate escalates only if I_t^e is not decision-sufficient; the fog/cloud agent uses the DT context c_t ; schedule/urgent actions need human approval.

conformal RUL intervals exist for bearings [12], but without a decision or routing layer. Conformal risk control [13] bounds expected monotone losses, and conformal alignment certifies edge–cloud cascades for classification [7].

Gap. To the best of our audit of 42 papers (Table I), no study uses conformal RUL intervals computed on the edge as the escalation rule of a hierarchical PdM system, relates escalation to maintenance-decision boundaries, and evaluates it against the gates of existing systems on public run-to-failure data with different degradation regimes. The contribution is methodological and evaluative; the architecture is not new.

III. HERA-PdM METHODOLOGY

Fig. 1 shows the four layers. “Agents” are software components with fixed responsibilities and messages; no language model is used anywhere.

Edge Monitoring Agent. Each vibration snapshot is reduced to $x_t \in \mathbb{R}^{28}$ (per axis: RMS, kurtosis, skewness, peak-to-peak, crest, impulse and shape factors, spectral centroid, six log band energies), z-scored against the asset’s first snapshots and compressed by $\text{sign}(z) \log(1+|z|)$. The health index is $h_t = \log(1 + \frac{1}{d} \sum_j |z_{t,j}|)$; the anomaly score a_t is the smoothed h_t standardized by the asset’s raw baseline, with alarm threshold τ_a at the 99th percentile of healthy training snapshots. A tiny model maps a short window to three RUL quantiles.

Digital Twin State Agent. The twin is a synchronized state record $DT_t = \{\text{condition, sensor state, } h_t, \text{ context } c_t, \hat{R}_t,$

$L_t, u_t, \text{ anomaly state}\}$ updated from a 12-byte sync (t, h_t, a_t) per snapshot; $c_t \in \mathbb{R}^6$ holds fast/slow averages of h_t , their difference, running maximum, slope and cumulative excess over baseline. Absolute operating time is excluded because lifetimes vary by an order of magnitude.

RUL Prognostics Agent. A sequence model g_ϕ maps a longer window X_t and c_t to non-crossing quantiles $(\hat{q}_t^{\alpha/2}, \hat{q}_t^{0.5}, \hat{q}_t^{1-\alpha/2})$ trained with the pinball loss on $y_t = \min(\text{RUL}_t, R_{\max})$.

Calibration (CQR-CV+). The non-test assets are split into $K=4$ condition-balanced groups; inner model $g^{(-k)}$ is trained without group k , which yields out-of-fold scores $s_i = \max(\hat{q}_i^{\alpha/2} - y_i, y_i - \hat{q}_i^{1-\alpha/2})$ for every non-test snapshot i in group $k(i)$. Following CV+ [11] with CQR scores [3],

$$L_t = Q_\alpha^- \{\hat{q}_t^{\alpha/2, (-k(i))} - s_i\}, \quad (1)$$

$$U_t = Q_{1-\alpha}^+ \{\hat{q}_t^{1-\alpha/2, (-k(i))} + s_i\}, \quad (2)$$

where Q_α^- and $Q_{1-\alpha}^+$ are the $\lfloor \alpha(n+1) \rfloor$ -th and $\lceil (1-\alpha)(n+1) \rceil$ -th smallest values over the n calibration snapshots; $\hat{R}_t$ is the mean inner median. CV+ guarantees coverage $\geq 1 - 2\alpha$ (typically $\approx 1 - \alpha$) if assets are exchangeable; snapshots within an asset are dependent, so the guarantee is approximate and we test it. Split CQR on three held-out bearings is the baseline calibrator.

Decision and escalation. With critical and planning horizons H_c and H_p , the action of a value v is $r(v) = \text{URGENT}$ if $v \leq H_c$, SCHEDULE if $H_c < v \leq H_p$, else CONTINUE; the Maintenance Decision Agent acts on $\pi_t = r(L_t)$ (CONTINUE becomes INSPECT if $a_t > \tau_a$). SCHEDULE/URGENT are recommendations with an explanation card (RUL, interval, anomaly evidence, top-3 SHAP [14] contributors, rationale) that require approval. With $B = \{H_c, H_p\}$, the edge escalates if its interval is ambiguous or contradicts anomaly evidence:

$$e_t = \mathbb{I}[\exists h \in B: L_t^e \leq h < U_t^e] \vee \mathbb{I}[a_t > \tau_a \wedge r(L_t^e) = C], \quad (3)$$

where C denotes CONTINUE. A local CONTINUE needs $L_t^e > H_p \geq H_c$, so a critical asset ($y_t \leq H_c$) is left running locally only if $y_t \notin I_t^e$:

$$\Pr(e_t=0, \pi_t=\text{CONTINUE}, y_t \leq H_c) \leq \Pr(y_t \notin I_t^e), \quad (4)$$

and false local URGENT actions are bounded likewise. Unlike a width threshold, the gate ignores wide intervals inside one action region and escalates narrow ones crossing a boundary.

Cost model. Payloads are counted in float32: always-cloud streams x_t (112 B/snapshot); HERA sends the 12 B sync plus the untransmitted part of the window when $e_t=1$. Latency is edge time + $e_t(\text{RTT} + \text{cloud time})$ with an assumed RTT of 50 ms.

IV. EXPERIMENTAL SETUP

Data. FEMTO-ST/PRONOSTIA [9]: 17 bearings under three conditions, 0.1 s snapshots every 10 s (24,889 snapshots), full test trajectories; several bearings fail abruptly. XJTU-SY [10]: 15 bearings under 35 Hz/12 kN, 37.5 Hz/11 kN and 40 Hz/10 kN, 1.28 s snapshots every

TABLE II

RUL ACCURACY AND CALIBRATION ON HELD-OUT BEARINGS (MEAN $\pm$ STD OVER 3 SEEDS; NOMINAL COVERAGE 0.90). RAW: UNCALIBRATED QUANTILES; SPLIT: CQR CALIBRATED ON 3 HELD-OUT BEARINGS PER FOLD; CV+: CQR-CV+ CALIBRATED ON ALL NON-TEST BEARINGS. MARKS: ARCHITECTURE SELECTED ON INNER-VALIDATION PINBALL LOSS FOR FEMTO ($\dagger$) AND XJTU-SY (*).

Model	FEMTO/PRONOSTIA (17 bearings, $R_{\max}=3000$ s)					XJTU-SY (15 bearings, $R_{\max}=6000$ s)				
	RMSE (s)	MAE (s)	Raw cov.	Cov.	MPIW (s)	RMSE (s)	MAE (s)	Raw cov.	Cov.	MPIW (s)
Edge-CNN* (split)	1114 $\pm$ 40	851 $\pm$ 19	0.37	0.88 $\pm$ 0.02	2080	2278 $\pm$ 389	1859 $\pm$ 330	0.65	0.97 $\pm$ 0.00	3716
Edge-GRU $\dagger$ (split)	820 $\pm$ 34	518 $\pm$ 62	0.71	0.84 $\pm$ 0.04	1877	2531 $\pm$ 642	2004 $\pm$ 601	0.85	0.87 $\pm$ 0.08	4184
LSTM (split)	1208 $\pm$ 118	772 $\pm$ 102	0.50	0.85 $\pm$ 0.04	2263	1516 $\pm$ 127	886 $\pm$ 104	0.90	0.91 $\pm$ 0.01	2436
LSTM+DT $\dagger$ (split)	1124 $\pm$ 48	751 $\pm$ 54	0.37	0.84 $\pm$ 0.04	2154	1698 $\pm$ 58	1130 $\pm$ 68	0.89	0.88 $\pm$ 0.03	3107
GRU (split)	1171 $\pm$ 114	751 $\pm$ 84	0.40	0.86 $\pm$ 0.03	2151	1516 $\pm$ 39	1027 $\pm$ 73	0.89	0.89 $\pm$ 0.02	2404
GRU+DT* (split)	1140 $\pm$ 89	771 $\pm$ 66	0.32	0.89 $\pm$ 0.03	2286	1647 $\pm$ 94	1226 $\pm$ 153	0.86	0.83 $\pm$ 0.06	3042
TCN (split)	1159 $\pm$ 70	827 $\pm$ 72	0.16	0.85 $\pm$ 0.02	2241	1666 $\pm$ 119	1303 $\pm$ 145	0.52	0.95 $\pm$ 0.01	3691
TCN+DT (split)	1153 $\pm$ 136	855 $\pm$ 115	0.16	0.89 $\pm$ 0.01	2456	1496 $\pm$ 118	1100 $\pm$ 138	0.62	0.81 $\pm$ 0.11	3477
CNN-LSTM (split)	1106 $\pm$ 74	674 $\pm$ 34	0.40	0.85 $\pm$ 0.02	2119	1335 $\pm$ 3	772 $\pm$ 21	0.85	0.90 $\pm$ 0.01	2416
CNN-LSTM+DT (split)	1108 $\pm$ 39	747 $\pm$ 54	0.28	0.89 $\pm$ 0.01	2355	1446 $\pm$ 30	895 $\pm$ 33	0.84	0.91 $\pm$ 0.00	2686
Selected edge (CV+)	753 $\pm$ 5	501 $\pm$ 22	0.75	0.94 $\pm$ 0.01	2193	1938 $\pm$ 251	1677 $\pm$ 232	0.69	0.99 $\pm$ 0.01	4134
Selected cloud+DT (CV+)	892 $\pm$ 27	665 $\pm$ 14	0.28	0.96 $\pm$ 0.02	2506	1422 $\pm$ 38	1077 $\pm$ 59	0.91	0.93 $\pm$ 0.01	3030
Selected cloud, no DT (CV+)	954 $\pm$ 16	653 $\pm$ 13	0.35	0.95 $\pm$ 0.02	2551	1305 $\pm$ 54	757 $\pm$ 66	0.90	0.94 $\pm$ 0.01	2337

TABLE III

HIERARCHICAL POLICIES, ABLATIONS AND SYSTEM COST (MEAN $\pm$ STD OVER 3 SEEDS, ALL HELD-OUT BEARINGS). ESC.: ESCALATED SNAPSHOTS; UNSAFE: CRITICAL SNAPSHOTS ($y \leq H_c$) LEFT WITHOUT A MAINTENANCE ACTION; FM: FALSE MAINTENANCE ON HEALTHY SNAPSHOTS ($y > H_p$); B/s: PAYLOAD BYTES PER SNAPSHOT (ANALYTIC). $\dagger$ ESCALATION BUDGET MATCHED TO HERA ON CALIBRATION DATA.

Policy	FEMTO/PRONOSTIA					XJTU-SY				
	Esc. (%)	Unsafe (%)	FM (%)	RMSE (s)	B/s	Esc. (%)	Unsafe (%)	FM (%)	RMSE (s)	B/s
Always-edge	0.0	0.6 $\pm$ 1.1	96.9 $\pm$ 1.4	752	0	0.0	0.0 $\pm$ 0.0	77.1 $\pm$ 7.1	1485	0
Always-cloud	100.0	0.0 $\pm$ 0.0	96.9 $\pm$ 2.7	891	112	100.0	9.2 $\pm$ 1.7	69.1 $\pm$ 6.1	1421	112
Cloud point + fixed thr.	100.0	49.5 $\pm$ 2.7	21.7 $\pm$ 6.8	1090	112	100.0	34.2 $\pm$ 9.2	8.3 $\pm$ 2.1	1689	112
HERA-full	97.1	0.6 $\pm$ 1.1	95.0 $\pm$ 0.5	892	123	79.1	9.2 $\pm$ 1.7	53.2 $\pm$ 13.3	1381	104
HERA w/o anomaly term	97.1	0.6 $\pm$ 1.1	95.0 $\pm$ 0.5	892	123	79.1	9.2 $\pm$ 1.7	53.2 $\pm$ 13.3	1381	104
HERA-no-DT	97.1	0.6 $\pm$ 1.1	96.9 $\pm$ 1.4	953	123	79.1	12.1 $\pm$ 4.3	25.5 $\pm$ 7.7	1303	104
HERA-no-UQ (point margin)	48.1	56.2 $\pm$ 5.3	15.1 $\pm$ 2.8	989	89	67.8	34.2 $\pm$ 9.2	8.0 $\pm$ 2.2	2062	91
HERA-fixed (anomaly trig.)	3.9	0.6 $\pm$ 1.1	96.9 $\pm$ 1.4	760	27	5.4	0.0 $\pm$ 0.0	77.1 $\pm$ 7.1	1541	20
Width gate $\dagger$	97.2	0.6 $\pm$ 1.1	95.1 $\pm$ 0.5	892	123	79.2	9.2 $\pm$ 1.7	53.3 $\pm$ 13.3	1381	104
Linear score gate $\dagger$	94.3	0.6 $\pm$ 1.1	95.7 $\pm$ 0.9	892	121	77.9	9.2 $\pm$ 1.7	53.8 $\pm$ 13.3	1378	103
Random gate $\dagger$	96.4	0.0 $\pm$ 0.0	96.8 $\pm$ 2.6	887	124	71.6	6.5 $\pm$ 1.4	71.4 $\pm$ 6.3	1441	124

minute (9,216 snapshots), mostly gradual degradation. RUL caps are $R_{\max}=3000$ s (FEMTO) and 6000 s (XJTU-SY); $H_c=0.2R_{\max}$, $H_p=0.6R_{\max}$; windows are 16/64 and 8/32 snapshots (edge/cloud). SCANIA Component X [15] (dataset DOI 10.5878/bnh5-ka77, CC BY 4.0): 23,550 trucks with irregular operational readouts (counters and histograms), 2,272 with an in-study repair and the rest right-censored; we use it only for unit-level alarm timing (Sec. V).

Protocol. Condition-balanced bearing-grouped folds (4 for FEMTO, 5 for XJTU-SY) test every bearing once. Each architecture’s hyper-parameters (width 32–128, 8 or 25 epochs, input noise) and the edge/cloud architecture choice were selected on inner validation bearings from training bearings only (FEMTO: Edge-GRU/LSTM+DT; XJTU-SY: Edge-CNN/GRU+DT). Three seeds per model; metrics are pooled over test bearings and reported as mean $\pm$ std; paired Wilcoxon tests use bearings as units.

Baselines and metrics. LSTM, GRU, TCN [8] and CNN-LSTM with and without DT context, MSE point mod-

els; always-edge, always-cloud and point-with-fixed-threshold policies; ablations without uncertainty (point margin), DT or anomaly term; anomaly-trigger (as in [5]), width (confidence-style deferral), linear-score and random gates at HERA’s budget. We report RMSE, MAE, coverage (Cov.), mean interval width (MPIW), *unsafe rate* (critical snapshots with CONTINUE/INSPECT), false maintenance (FM, SCHEDULE/URGENT when $y_t > H_p$), escalation, payload, latency (single-thread host CPU, Apple M2 Pro; no energy measured) and, per bearing, the timing of the first sustained maintenance alarm.

V. RESULTS AND DISCUSSION

RQ1: accuracy. Table II shows opposite regimes. On FEMTO, the 2.5k-parameter edge model was most accurate (RMSE 753 s with CQR-CV+ vs. 892 s for LSTM+DT), because abrupt failures leave larger models nothing to exploit. On XJTU-SY, fog/cloud models were clearly better (best split-trained 1335 s, CNN-LSTM) than the selected edge model (2278 s). DT context lowered RMSE for 3 of 4 architectures




Fig. 2. Fog/cloud median RUL and 90% CQR-CV+ interval before failure (seed 0; per dataset the test bearing with most degradation-phase snapshots). Dotted line: H_c .

TABLE IV

TOP: OUTCOME OF EACH BEARING'S FIRST SUSTAINED (3 SNAPSHOTS) SCHEDULE/URGENT RECOMMENDATION (NUMBER OF BEARINGS, MEAN OVER 3 SEEDS): T TIMELY ($H_c \leq \text{RUL} \leq R_{\max}$), L LATE ($\text{RUL} < H_c$ OR NEVER), P PREMATURE ($\text{RUL} > R_{\max}$); LEAD: MEDIAN RUL AT TIMELY ALARMS (MIN). BOTTOM: SENSITIVITY OF UNSAFE/FM (%) FOR HERA-FULL AND THE POINT POLICY TO THE HORIZONS AND TO $R_{\max}$.

Policy	FEMTO (17 bearings)				XJTU-SY (15 bearings)			
	T	L	P	lead	T	L	P	lead
Always-edge	1.0	0.0	16.0	38	2.0	0.0	13.0	44
Always-cloud	1.0	0.0	16.0	38	1.0	1.0	13.0	49
Point + thr.	1.0	5.3	10.7	40	5.0	4.0	6.0	54
HERA-full	1.0	0.0	16.0	38	1.0	1.0	13.0	49
Width gate	1.0	0.0	16.0	38	1.0	1.0	13.0	49
Setting	HERA		Point		HERA		Point	
$H_c, H_p=0.1, 0.4R_{\max}$	13/68		54/15		16/15		50/5	
$H_c, H_p=0.2, 0.6R_{\max}$	1/95		49/22		9/53		34/8	
$H_c, H_p=0.3, 0.8R_{\max}$	0/100		37/39		2/81		23/42	
$R_{\max} \times \frac{1}{2}$ (3 seeds)	22/31		35/19		7/22		56/6	
$R_{\max} \times 2$ (3 seeds)	0/98		24/44		0/81		23/31	

on FEMTO but only 1 on XJTU-SY; cross-fitted ensembles (CV+) also improved point accuracy.

RQ2: calibration. Split CQR with three calibration bearings under-covered on FEMTO (0.84–0.89). CQR-CV+, which calibrates on all non-test bearings, reached 0.94 (edge) and 0.96 (cloud) on FEMTO and 0.93–0.99 on XJTU-SY, at the price of wider intervals (MPIW 2193 s and 3030 s). Per-bearing coverage still varied strongly (lowest cloud coverage 0.30 and 0.07), so exchangeability between bearings remains the binding assumption. Fig. 2 shows intervals narrowing once degradation is visible, far earlier on XJTU-SY.

RQ4: unsafe decisions. Point estimates with fixed thresholds left 49.5% (FEMTO) and 34.2% (XJTU-SY) of critical snapshots without a maintenance action, versus 0.6% and 9.2% for HERA (Table III); the differences were -48.8 pp, 95% CI $[-65.4, -32.0]$ (FEMTO) and -25.0 pp, 95% CI $[-41.8, -10.7]$ (XJTU-SY) over bearings (Wilcoxon $p < 0.001$, $p = 0.004$), and false maintenance rose by 73.3 pp,

95% CI $[60.2, 84.7]$ and 44.9 pp, 95% CI $[17.0, 60.8]$. Sweeping α or a point margin δ (Fig. 3a,c) shows that calibrated and margin-adjusted point policies lie close to one frontier: calibration contributes a principled, held-out choice of operating point rather than better attainable decisions.

RQ3 and RQ5: escalation and system trade-off. On FEMTO, 90% edge intervals (2193 s on a 3000 s range) almost always straddled H_p , so HERA escalated 97.1% of snapshots, sent more payload than always-cloud streaming (123.0 vs. 112.0 B) and gave no decision benefit. On XJTU-SY, HERA escalated 79.1%, matched always-cloud safety (9.2% vs. 9.2%) with lower false maintenance (53.2% vs. 69.1%), better RMSE (1381 vs. 1421 s), 104.5 instead of 112.0 B/snapshot and mean latency 40.1 vs. 50.6 ms. Two caveats temper this: the FM difference to always-cloud (-15.9 pp, 95% CI $[-35.2, 0.0]$) comes from the two long-lived condition-3 bearings and its interval reaches zero, and the edge-only calibrated policy traced a better frontier than HERA (Fig. 3c) because the conservative edge intervals rarely under-predict. At matched budget, the width gate gave identical results (9.2%/53.3%) and the linear-score gate nearly so (Fig. 3b,d); the anomaly trigger escalated only 5.4% (anomaly AUROC 0.81/0.80). Relaxing α_e to 0.3 reduced escalation to 83.0%/67.9% while unsafe decisions rose to 7.7%/9.2%. Removing the DT context lowered FM on XJTU-SY (25.5%) at higher unsafe rate (12.1%), so the DT context is not uniformly beneficial.

Per-bearing timing and sensitivity. Table IV exposes the main practical limitation: because the 90% lower bound falls below H_p early, every calibrated policy raised its first sustained maintenance alarm while the true RUL still exceeded $R_{\max}$ for 16.0 of 17 FEMTO and 13.0 of 15 XJTU-SY bearings, whereas the point policy was late for 5.3 and 4.0 bearings. Neither is acceptable for scheduling; per-snapshot marginal coverage is the wrong target for alarm timing. In all five horizon settings and all four halved/doubled- $R_{\max}$ runs, HERA left fewer critical snapshots without action than the point policy; with a halved cap on XJTU-SY, HERA reached low false maintenance at the always-cloud unsafe rate (Table IV).

Bearing-level risk control. We also calibrated the alarm offset directly for timing: an onset-gated two-stage alarm whose offset λ was chosen by conformal risk control [13] on the CV+ out-of-fold predictions of the non-test bearings, so that $\Pr(\text{late or missing alarm}) \leq \varepsilon$ per bearing. At $\varepsilon = 0.1$ the held-out late rate stayed within the target (2.0% FEMTO, 4.4% XJTU-SY vs. 35.3% and 26.7% for the point policy), but only by alarming early: 94.1% and 80.0% of alarms were premature and the onset gate escalated 99.8% and 98.7% of snapshots. The reason is a heavy tail of required offsets: the median bearing needed an offset of 0.24 (FEMTO) and 0.20 (XJTU-SY) of $R_{\max}$ to be warned in time, but the two least predictable bearings needed 0.67–0.75 and 0.58–0.74. This is a sample-size effect:

Proposition 1. Let a fraction π of n calibration units be late for every admissible offset (no usable warning). CRC with the binary lateness loss returns a non-trivial offset at level ε only




Fig. 3. (a,c) Unsafe decisions vs. false maintenance when sweeping α (calibrated; circles at $\alpha=0.1$) or a margin δ (point). (b,d) Unsafe decisions vs. escalation budget for gates on the same $\alpha=0.1$ edge intervals; green line: HERA with α_e swept.

TABLE V

FLEET SCALE (SCANIA COMPONENT X, 23,550 TRUCKS, 2,272 FAILING; 50 RANDOM CALIBRATION/TEST SPLITS). LATE: FAILING TRUCKS WITHOUT ALARM BEFORE $RUL < H_c$ (MEAN [95TH PCT.]); FA: HEALTHY TRUCKS ALARMED $> H_p$ EARLY; ESC.: READOUTS SENT TO THE CLOUD.

Policy	n_{cal}	Late (%)	FA (%)	Esc. (%)
Edge, fixed $p \geq 0.5$	—	22.5	52.1	0
Cloud, fixed $p \geq 0.5$	—	40.5	28.6	100
Cloud, CRC $\varepsilon=0.1$	12	7.7 [19.8]	73.9	100
Cloud, CRC $\varepsilon=0.1$	1000	9.8 [12.2]	66.2	100
Edge, CRC $\varepsilon=0.1$	1000	9.9 [11.8]	68.1	0
HERA , CRC $\varepsilon=0.1$	1000	9.9 [12.2]	66.1	73
HERA , CRC $\varepsilon=0.05$	1000	4.9 [6.7]	77.1	79
HERA , CRC $\varepsilon=0.2$	1000	19.8 [23.3]	50.1	64

TABLE VI

TOP: DETECTION OF IMMINENT FAILURE ($RUL \leq H_c$) PER SNAPSHOT, THRESHOLDS SET ON TRAINING DATA (MEAN OVER 3 SEEDS). BOTTOM: COMPUTE OF THE SELECTED EDGE AND FOG/CLOUD MODELS (HOST CPU, 1 THREAD, BATCH 1).

Detector	FEMTO			XJTU-SY		
	Prec.	Rec.	AUROC	Prec.	Rec.	AUROC
Anomaly a_t	0.27	0.26	0.81	0.21	0.34	0.80
Isolation Forest	0.27	0.15	0.77	0.15	0.13	0.81
Edge $L_t^e \leq H_c$	0.07	0.68	0.70	0.06	0.94	0.82
Cloud $L_t^c \leq H_c$	0.05	0.82	0.65	0.17	0.76	0.81
Model	Params	MACs	ms	Params	MACs	ms
Edge model	2.5k	0.03M	0.15	2.7k	0.02M	0.07
Cloud model	62.1k	3.61M	1.11	13.6k	0.38M	0.62

if $(n\pi + 1)/(n + 1) \leq \varepsilon$, i.e. $\pi < \varepsilon$ and $n \geq (1 - \varepsilon)/(\varepsilon - \pi)$.

Proof. The empirical risk of every offset is at least π , and CRC requires $(n\hat{R} + 1)/(n + 1) \leq \varepsilon$ [13]. $\square$

For $\varepsilon=0.1$ and $\pi=0.05$ this needs $n \geq 18$, beyond the ≈ 12 bearings per fold available here; below it CRC must alarm at the earliest possible offset.

Fleet-scale validation. To test Proposition 1 where n is large, we applied the same unit-level CRC to 23,550 Scania trucks: a logistic-regression edge model and a gradient-boosted cloud model (vehicle-grouped 5-fold cross-fitting, AUROC 0.83 and 0.83) score $\Pr(RUL \leq H_p)$ per readout with $H_c=12$ and $H_p=48$ time units; alarms need two consecutive readouts; HERA opens a latched edge gate at half the cloud threshold. Table V averages 50 random calibration/test splits. Fixed thresholds were late for 22.5% (edge) and 40.5% (cloud) of failing trucks, whereas CRC hit its target (9.8% at $\varepsilon=0.1$, 4.9% at 0.05, 19.8% at 0.2). Growing the calibration fleet from 12 to 1000 failing trucks tightened the 95th-percentile late rate from 19.8% to 12.2% and lowered false alarms from 73.9% to 66.2%. HERA matched always-cloud decisions (9.9%/66.1% vs. 9.8%/66.2%) while escalating only 73% of readouts; the

edge model alone was almost as good (68.1% false alarms) because the cloud model was not more accurate. False alarms stayed high because the failure signal is weak.

Detection and compute. Table VI shows why no simple trigger suffices: the edge anomaly score reached AUROC 0.81 and 0.80 but recall of only 0.26 and 0.34 at its training-calibrated threshold, whereas the calibrated lower bounds recalled 0.68–0.94 of critical snapshots at low precision, the safety/false-alarm trade-off seen in Fig. 3. The edge models need tens of thousands of MACs and well under a millisecond; the fog/cloud models need roughly 20–120 times more MACs. Run unchanged in an ARM64 Linux container capped at one CPU and 1 GB (Apple M2 Pro host, not embedded hardware), the exported models reproduced host outputs to below 10^{-6} and took at most 0.65 ms (edge) and 2.6 ms (cloud) per window, plus 2.8 ms for feature extraction, with 261 MB peak memory; with the container pinned to one core, p99 stayed below 2.5 ms.

Explainability and oversight. Expected-gradient SHAP on 300 degradation-phase snapshots ranked low-frequency band energies and the spectral centroid (FEMTO) and the DT

TABLE VII
DEPLOYMENT GUIDE DERIVED FROM THE MEASURED RESULTS (FEMTO / XJTU-SY).

Plant condition	Recommended policy	Evidence
Safety-critical, any fleet	Decide on calibrated lower bound	unsafe 0.6%/9.2% vs. 49.5%/34.2% for point thresholds
Degradation visible early	Uncertainty-gated escalation	XJTU-SY: same unsafe rate as always-cloud, FM 53.2% vs. 69.1% (CI reaches 0)
Abrupt, warning-free failures	Always-edge calibrated	FEMTO: HERA escalated 97.1%, no benefit; edge RMSE 753 s best
Limited bandwidth	Gate on interval width	width gate equal to decision-sufficiency gate on both datasets
Alarm-timing guarantee needed	Conformal risk control with $n \geq (1-\epsilon)/(\epsilon-\pi)$ units	bearings (≈ 12 units): premature 94.1%/80.0%; trucks (1000): late 9.9% at target 10%
Limited operator capacity	Relax α or add onset detection	approval load 95.5%/55.8% at $\alpha=0.1$

fast health-index average and mid- and high-frequency band energies (XJTU-SY) highest; additivity errors were 16% and 3% of the explained deviation. An automatically generated card (FEMTO) reads:

Asset Bearing1_3, snapshot 2078. **RUL** 578 s, 90% interval [0, 2138] s. **Anomaly** $a_t=15.8$ (threshold 17.04, no alarm). **Top SHAP** h_band0_1k (-1489 s), v_band2_4k (-1244 s), $h_centroid$ (+1175 s). **Recommendation** urgent maintenance — *requires human approval*. **Rationale** lower bound $0s \leq H_c$.

All fields come from the predictors via a fixed template. Approval requests covered 95.5% (FEMTO) and 55.8% (XJTU-SY) of HERA snapshots, an operator load that shows calibrated but wide intervals shift burden to humans.

Lessons for deployment. Table VII condenses the evidence. (i) *Calibrate before routing*: calibrated intervals removed most unsafe point-threshold decisions on both datasets. (ii) *Hierarchy pays only if degradation is observable early*: escalation helped on XJTU-SY and not on FEMTO. (iii) *Small fleets strain calibration*: three calibration bearings under-covered, cross-fitting on all assets fixed pooled coverage, and per-asset coverage still varied. (iv) *The interval width matters more than the gate*: decision-sufficiency and width gates coincided. (v) *Unit-level guarantees need fleets*: with a dozen units risk control alarms early (Proposition 1); with a thousand trucks it met its target. (vi) *Wide intervals move work to humans*: approval load, not model accuracy, limits human-centric use. Limitations: 32 bearings, approximate exchangeability, a state-based DT, no timing on embedded hardware, analytic communication, and no deployment or user study.

VI. CONCLUSION

Cross-fitted conformal RUL intervals removed most unsafe decisions of point-threshold maintenance on two bearing datasets. Using them to gate edge-to-cloud escalation paid

off on gradually degrading bearings and on a truck fleet, where HERA matched always-cloud safety with a fraction of the cloud traffic, but not under abrupt failures; the specific decision-sufficiency rule was no better than a width gate. Calibrating alarm timing with conformal risk control is feasible only with enough units (Proposition 1): with a dozen bearings it alarmed early, whereas on 23,550 trucks it met its lateness target, at a false-alarm cost set by the weak failure signal. Better early-warning features are therefore the next lever, not routing rules.

REFERENCES

- [1] Y. Lei, N. Li, L. Guo, N. Li, T. Yan, and J. Lin, “Machinery health prognostics: A systematic review from data acquisition to RUL prediction,” *Mechanical Systems and Signal Processing*, vol. 104, pp. 799–834, 2018.
- [2] L. Ren, Y. Liu, X. Wang, J. Lü, and M. J. Deen, “Cloud-edge-based lightweight temporal convolutional networks for remaining useful life prediction in IIoT,” *IEEE Internet of Things Journal*, vol. 8, no. 16, pp. 12 578–12 587, 2021.
- [3] Y. Romano, E. Patterson, and E. J. Candès, “Conformalized quantile regression,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [4] X. Xu, Y. Lu, B. Vogel-Heuser, and L. Wang, “Industry 4.0 and industry 5.0—inception, conception and perception,” *Journal of Manufacturing Systems*, vol. 61, pp. 530–535, 2021.
- [5] R. B. Saleh, K. P. Dinh, B. Villanyi, and T.-S. Hy, “HAMA: A hierarchical adaptive multi-agent architecture for industrial IoT predictive maintenance,” *IEEE Access*, vol. 14, pp. 142 595–142 613, 2026.
- [6] Z. Yao, Y. Xu, T. Xia, Z. Chen, E. Pan, and L. Xi, “EC-Distill-ZeroDiag: a cloud-edge collaborative framework for zero-shot industrial fault diagnosis via large language model distillation,” *Journal of Intelligent Manufacturing*, 2026.
- [7] J. Huang, S. Park, N. Paoletti, and O. Simeone, “Reliable inference in edge-cloud model cascades via conformal alignment,” arXiv preprint arXiv:2510.17543, 2025.
- [8] S. Bai, J. Z. Kolter, and V. Koltun, “An empirical evaluation of generic convolutional and recurrent networks for sequence modeling,” arXiv preprint arXiv:1803.01271, 2018.
- [9] P. Nectoux, R. Gouriveau, K. Medjaher, E. Ramasso, B. Chebel-Morello, N. Zerhouni, and C. Varnier, “PRONOSTIA: An experimental platform for bearings accelerated degradation tests,” in *Proc. IEEE Int. Conf. Prognostics and Health Management (PHM’12)*, Denver, CO, USA, 2012, pp. 1–8, hAL: hal-00719503.
- [10] B. Wang, Y. Lei, N. Li, and N. Li, “A hybrid prognostics approach for estimating remaining useful life of rolling element bearings,” *IEEE Transactions on Reliability*, vol. 69, no. 1, pp. 401–412, 2020.
- [11] R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani, “Predictive inference with the jackknife+,” *The Annals of Statistics*, vol. 49, no. 1, 2021.
- [12] S. Wang, Y. Vidal, and F. Pozo, “Uncertainty-aware bearing remaining useful life prediction based on conformal prediction,” in *Proc. PHM Society European Conference*, vol. 9, no. 1, 2026, pp. 1–10.
- [13] A. N. Angelopoulos, S. Bates, A. Fisch, L. Lei, and T. Schuster, “Conformal risk control,” in *International Conference on Learning Representations (ICLR)*, 2024.
- [14] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017, pp. 4765–4774.
- [15] Z. Kharazian, T. Lindgren, S. Magnússon, O. Steinert, and O. Andersson Reyna, “SCANIA Component X dataset: a real-world multivariate time series dataset for predictive maintenance,” *Scientific Data*, vol. 12, 2025.